

Lecture 15: RAI and Smart Pointers

CS 106L, Fall '20

Agenda

- Exceptions
- RAII
- Smart pointers

How many code paths are in this function?

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 1 - favors neither chocolate nor milkshakes

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 1 - favors neither chocolate nor milkshakes

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 1 - favors neither chocolate nor milkshakes

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 1 - favors neither chocolate nor milkshakes

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 2 - favors milkshakes

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```


Code Path 2 - favors milkshakes

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 2 - favors milkshakes

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 2 - favors milkshakes

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 2 - favors milkshakes

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 3 - favors chocolate (and possibly milkshakes)

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 3 - favors chocolate (and possibly milkshakes)

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 3 - favors chocolate (and possibly milkshakes)

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Code Path 3 - favors chocolate (and possibly milkshakes)

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```


Are there any more code paths?

```
string evaluate_sweet_tooth_and_return_name( Person p ) {  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
    return p.first() + " " + p.last();  
}
```

Hidden Code Paths

There are (at least) 23 code paths in the code before!

- 1 – Copy constructor of Person parameter, may throw.
- 5 – Constructor of temp string, may throw.
- 6 – Call to favorite_food, favorite_drink, first (2), last (2), may throw.
- 10 – Operators may be user-overloaded, may throw.
- 1 – Copy constructor of string for return value, may throw.

Brief aside: exceptions

Aside: Exceptions

Exceptions are a way to transfer control and information to a (potential) exception handler

```
try {  
    // code associated with exception handler  
} catch ( [exception type] e ) {  
    // exception handler  
} catch ( [exception type] e ) {  
    // exception handler  
} // etc.
```

What might go wrong here? (Answer in chat)

```
string evaluate_sweet_tooth_and_return_name( int id_number ) {  
    Person *p = new Person(id_number);  
  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
  
    auto result = p.first() + " " + p.last();  
    delete p;  
  
    return result;  
}
```

Can we guarantee this function won't leak memory?

```
string evaluate_sweet_tooth_and_return_name( int id_number ) {  
    Person *p = new Person(id_number);  
  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
  
    auto result = p.first() + " " + p.last();  
    delete p;  
  
    return result;  
}
```

The “delete” won’t happen if there’s an exception first!

```
string evaluate_sweet_tooth_and_return_name( int id_number ) {  
    Person *p = new Person(id_number);  
  
    if ( p.favorite_food() == "chocolate" ||  
        p.favorite_drink() == "milkshake" ) {  
        cout << p.first() << " "  
             << p.last() << " has a sweet tooth" << endl;  
    }  
  
    auto result = p.first() + " " + p.last();  
    delete p;  
  
    return result;  
}
```

Lots of kinds of resources need to be released

Resources that need to be returned.

	Acquire	Release
• Heap memory	new	delete
• Files	open	close
• Locks	try_lock	unlock
• Sockets	socket	close

How do we guarantee resources get released, even if there are exceptions?

RAII
(Resource Acquisition is
Initialization)

RAII

"The best example of why I
shouldn't be in marketing"

"I didn't have a good day when I
named that"

-- Bjarne Stroustrup

 **Questions?** 

What is RAI?

- All resources should be **acquired** in the constructor.
- All resources should be **released** in the destructor.

What is RAI?

- All resources should be **acquired** in the constructor.
- All resources should be **released** in the destructor.

What's the rationale for this?

- There should never be a “half-valid” state of the object--object is usable immediately after creation.
- The destructor is always called (even with exceptions), so the resource is always freed.

You learned this in CS 106B. Is it RAll-compliant?

```
void printFile () {  
    ifstream input;  
    input.open("hamlet.txt");  
  
    string line;  
    while (getline(input, line)) { // might throw exception  
        cout << line << endl;  
    }  
  
    input.close();  
}
```

Nope - resource not acquired in ctor/released in dtor

```
void printFile () {  
    ifstream input;  
    input.open("hamlet.txt");  
  
    string line;  
    while (getline(input, line)) { // might throw exception  
        cout << line << endl;  
    }  
  
    input.close();  
}
```


This fixes it!

```
void printFile () {  
    ifstream input("hamlet.txt");  
  
    string line;  
    while (getline(input, line)) { // might throw exception  
        cout << line << endl;  
    }  
  
    // no close call needed!  
} // stream destructor, releases access to file
```

This is also not RAll-compliant!

```
void cleanDatabase (mutex& databaseLock,  
                    map<int, int>& database) {  
  
    databaseLock.lock();  
  
    // other threads will not modify database  
    // modify the database  
    // if exception thrown, mutex never unlocked!  
  
    databaseLock.unlock();  
}
```

This fixes it!

The `lock_guard` is an object whose sole job is to release the resource (unlock the mutex) when it goes out of scope

```
void cleanDatabase (mutex& databaseLock,  
                    map<int, int>& database) {  
  
    lock_guard<mutex> lock(databaseLock);  
  
    // other threads will not modify database  
    // modify the database  
    // if exception thrown, that's fine!  
  
    // no release call needed  
} // lock always unlocked when function exits.
```

How might lock_guard be implemented?



Here's a non-template version

```
class lock_guard {  
public:  
    lock_guard(mutex& lock) : acquired_lock(lock) {  
        acquired_lock.lock()  
    }  
    ~lock_guard() {  
        acquired_lock.unlock();  
    }  
private:  
    mutex& acquired_lock;  
}
```

RAII Summary

- Acquire resources in the constructor, release in the destructor.
- Clients of your class won't have to worry about mismanaged resources.

But what about RAI for memory?

This is where we're going!

R.11: Avoid calling `new` and `delete` explicitly

Reason

The pointer returned by `new` should belong to a resource handle (that can call `delete`). If the pointer returned by `new` is assigned to a plain/naked pointer, the object can be leaked.

Note

In a large program, a naked `delete` (that is a `delete` in application code, rather than part of code devoted to resource management) is a likely bug: if you have N `delete`s, how can you be certain that you don't need $N+1$ or $N-1$? The bug may be latent: it may emerge only during maintenance. If you have a naked `new`, you probably need a naked `delete` somewhere, so you probably have a bug.

Enforcement

(Simple) Warn on any explicit use of `new` and `delete`. Suggest using `make_unique` instead.

 **Questions?** 

Announcements

Assignment 2 has been released!

- For assignment 2, you'll be extending an existing implementation of a `HashMap` data structure to make it more similar to the `std::unordered_map`!
- Checkpoint 1 is due Thursday, November 12 at 11:59 PM
- Checkpoint 2 is due Wednesday, November 18 at 11:59 PM
- Avoid using late days for the first checkpoint, but feel free to use them for the second
- You may also work with a partner if you'd like (you'll only submit one assignment for the both of you)

Smart Pointers

(RAII for memory!)

We just saw how locks could be made RAll-compliant

```
void cleanDatabase (mutex& databaseLock,  
                    map<int, int>& database) {  
  
    databaseLock.lock();  
  
    // other threads will not modify database  
    // modify the database  
    // if exception thrown, mutex never unlocked!  
  
    databaseLock.unlock();  
}
```

... where the fix was to wrap it in a special object

```
void cleanDatabase (mutex& databaseLock,  
                    map<int, int>& database) {  
  
    lock_guard<mutex> lock(databaseLock);  
  
    // other threads will not modify database  
    // modify the database  
    // if exception thrown, that's fine!  
  
    // no release call needed  
} // lock always unlocked when function exits.
```

... so let's do it again!



You learned this in CS 106B -- is this RAll-compliant?


```
void rawPtrFn () {  
    Node* n = new Node;  
  
    // do some stuff with n..  
  
    delete n;  
}
```


You learned this in CS 106B -- is this RAII-compliant?

```
void rawPtrFn () {  
    Node* n = new Node;  
  
    // do some stuff with n...  
    // if exception thrown, n never deleted!  
    delete n;  
}
```

You learned this in CS 106B -- is this RAll-compliant?

```
void rawPtrFn () {  
    Node* n = new Node  
  
    // do some stuff w  
    // if exception th never deleted!  
    delete n;  
}
```



Solution: built-in “smart” (RAII-compliant) pointers

```
std::unique_ptr  
std::shared_ptr  
std::weak_ptr  
// std::auto_ptr is deprecated
```

`std::unique_ptr`

- Uniquely owns its resource and deletes it when the object is destroyed
- Cannot be copied (but can be moved, if non-const!)

std::unique_ptr

Before

```
void rawPtrFn () {  
    Node* n = new Node;  
    // do stuff with n...  
    delete n;  
}
```

After

```
void rawPtrFn () {  
    std::unique_ptr<Node> n(new Node);  
    // do some stuff with n  
} // Freed!
```

We can't copy an `std::unique_ptr`, but what if we could?

We can't copy an `std::unique_ptr`, but what if we could?


First we make a `unique_ptr`

```
unique_ptr<int> x;
```




We can't copy an `std::unique_ptr`, but what if we could?

We then make a copy of our `unique_ptr`



We can't copy an `std::unique_ptr`, but what if we could?


When `y` goes out of scope, it deletes the heap data



We can't copy an `std::unique_ptr`, but what if we could?

This leaves `x` pointing at deallocated data

```
unique_ptr<int> x;
```



We can't copy an `std::unique_ptr`, but what if we could?

If we dereference `x` or its destructor calls `delete`, we crash



We can't copy an `std::unique_ptr`, but what if we could?

If we dereference `x` or its destructor calls `delete`, we crash



We can't copy an `std::unique_ptr`, but what if we could?

- The `std::unique_ptr` class hence disallows copying!

 **Questions?** 

**But what if we wanted to have multiple
pointers to the same object?**

`std::shared_ptr`

- Resource can be stored by any number of `std::shared_ptr`s
- The resource is **deleted** when none of them points to it

std::shared_ptr

- Resource can be stored by any number of std::shared_ptrs
- The resource is **deleted** when none of them points to it

```
{
    std::shared_ptr<int> p1(new int);
    // Use p1
    {
        std::shared_ptr<int> p2 = p1;
        // Use p1 and p2
    }
    // Use p1
}
// The integer is deallocated!
```

std::shared_ptr

- Resource can be stored by any number of std::shared_ptrs
- The resource is **deleted** when none of them points to it

```
{  
    std::shared_ptr<int> p1(new int);  
    // Use p1  
    {  
        std::shared_ptr<int> p2 = p1;  
        // Use p1 and p2  
    }  
    // Use p1  
}  
// The integer is deallocated!
```

Important: This only works if new shared_ptrs are made by copying, not by constructing using the raw pointer!

How are `std::shared_ptrs` implemented?

How are these implemented? Reference counting!

- Idea: Store int that tracks how many `shared_ptrs` currently reference that data
 - Increment in copy constructor/copy assignment
 - Decrement in destructor or when overwritten with copy assignment
- Frees the resource when reference count hits 0

And finally, `std::weak_ptr`

Similar to `std::shared_ptr`, but doesn't contribute to the reference count.

To access the referenced data, must convert to `std::shared_ptr`

```
std::shared_ptr<T> tempSharedPtr = weakPtr.lock();
```

And finally, `std::weak_ptr`

Similar to `std::shared_ptr`, but doesn't contribute to the reference count.

To access the referenced data, must convert to `std::shared_ptr`

```
std::shared_ptr<T> tempSharedPtr = weakPtr.lock();
```

Used to observe a resource without influencing its lifetime (example use case: use `weak_ptr` to determine if a `shared_ptr` has deallocated a resource), or to break `shared_ptr` cycles.

Smart pointers: RAII wrapper for pointers

```
std::unique_ptr<T> up{new T};
```

```
std::shared_ptr<T> sp{new T};
```

```
std::weak_ptr<T> wp = sp;
```

But wait ... aren't we technically still using new?

```
std::unique_ptr<T> up{new T};
```

```
std::shared_ptr<T> sp{new T};
```

```
std::weak_ptr<T> wp = sp;
```

R.11: Avoid calling **new** and **delete** explicitly

There's another option!

```
std::unique_ptr<T> up{new T};  
std::unique_ptr<T> up = std::make_unique<T>();  
  
std::shared_ptr<T> sp{new T};  
std::shared_ptr<T> sp = std::make_shared<T>();  
  
std::weak_ptr<T> wp = sp;  
//can only be copy/move constructed (or empty)!
```


Which way is better?

```
std::unique_ptr<T> up{new T};  
std::unique_ptr<T> up = std::make_unique<T>();
```

Which way is better?

```
std::unique_ptr<T> up{new T};  
std::unique_ptr<T> up = std::make_unique<T>();
```

Answer:

Always use `std::make_unique<T>()`!

Final notes

- `std::unique_ptr`s are used often
- `std::shared_ptr` and `std::weak_ptr` are not used as often

 **Questions?** 

Live Code Demo:

Smart pointers in action

Next time

Guest lecture